



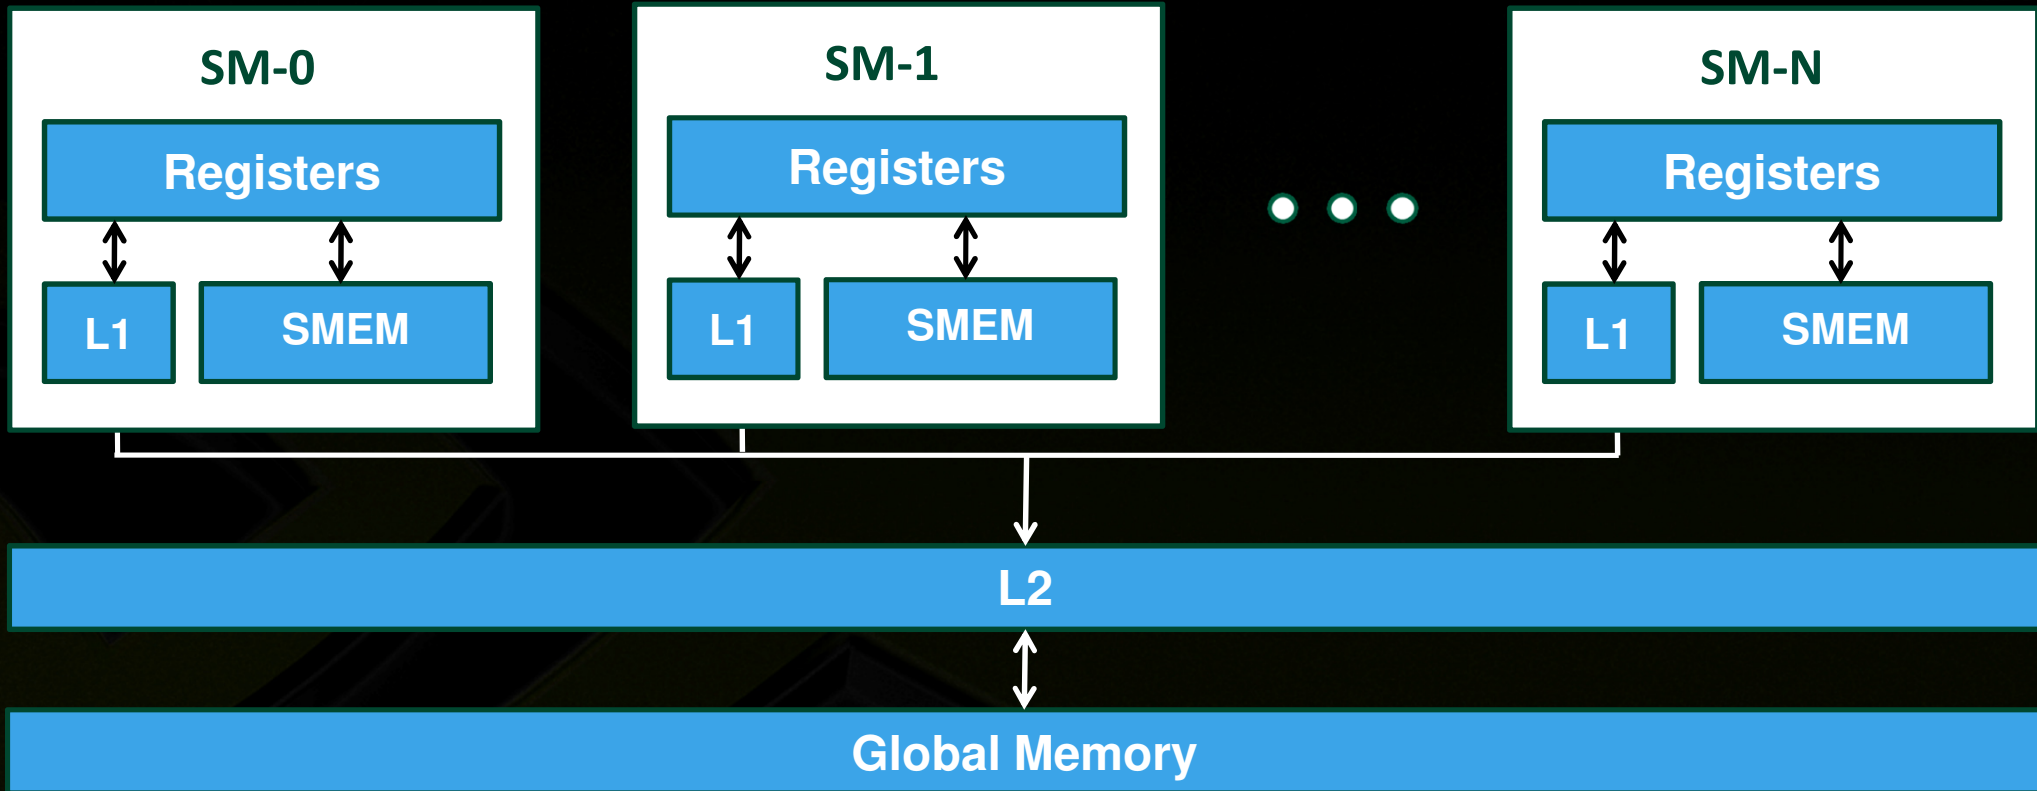
# **CUDA Warps and Occupancy**

**GPU Computing Webinar 7/12/2011**

**Dr. Justin Luitjens, NVIDIA Corporation**  
**Dr. Steven Rennich, NVIDIA Corporation**



# Fermi Architecture



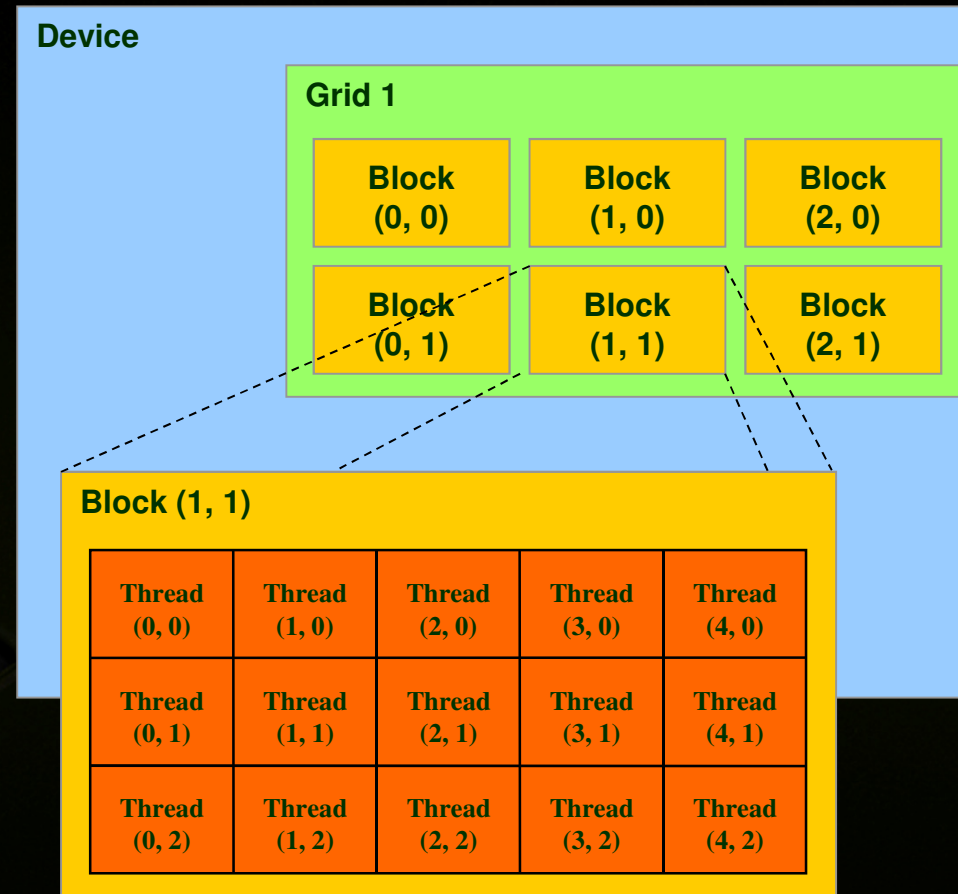
● **SM – Streaming multi-processors with multiple processing cores**

- Each SM contains 32 processing cores
- Execute in a Single Instruction Multiple Thread (SIMT) fashion
- Up to 16 SMs on a card for a maximum of 512 compute cores



# CUDA Programming Model Review

- A **grid** is composed of blocks which are completely independent
- A **block** is composed of threads which can communicate within their own block
- 32 threads form a **warp**
- Instructions are issued per warp
- If an operand is not ready the warp will stall
  - Context switch between warps when stalled
  - Context switch must be very fast




# Fast Context Switching

- Registers and shared memory are allocated for a block as long as that block is active
  - Once a block is active it will stay active until all threads in that block have completed
  - Context switching is very fast because registers and shared memory do not need to be saved and restored
- **Goal: Have enough transactions in flight to saturate the memory bus**
  - Latency can be hidden by having more transactions in flight
  - Increase active threads or Instruction Level Parallelism (ILP)
- Fermi can have up to 48 active warps per SM (1536 threads)

# Maximizing Memory Throughput



- **Increment of an array of 64M elements**
  - Two accesses per thread (load then store)
  - The two accesses are dependent, so really 1 access per thread at a time
- **Tesla C2050, ECC on, theoretical bandwidth: ~120 GB/s**



Several independent smaller accesses have the same effect as one larger one.

For example:

Four 32-bit  $\approx$  one 128-bit



# Occupancy



- **Occupancy = Active Warps / Maximum Active Warps**
- **Remember: resources are allocated for the entire block**
  - Resources are finite
  - Utilizing too many resources per thread may limit the occupancy
- **Potential occupancy limiters:**
  - Register usage
  - Shared memory usage
  - Block size

# Occupancy Limiters: Registers

- Register usage: compile with `--ptxas-options=-v`
- Fermi has 32K registers per SM
- Example 1
  - Kernel uses 20 registers per thread (+1 implicit)
  - Active threads =  $32K/21 = 1560$  threads
    - $> 1536$  thus an occupancy of 1
- Example 2
  - Kernel uses 63 registers per thread (+1 implicit)
  - Active threads =  $32K/64 = 512$  threads
  - $512/1536 = .3333$  occupancy
- Can control register usage with the nvcc flag: `--maxrregcount`

# Occupancy Limiters: Shared Memory

- Shared memory usage: compile with `--ptxas-options=-v`
  - Reports shared memory per block
- Fermi has either 16K or 48K shared memory
- Example 1, 48K shared memory
  - Kernel uses 32 bytes of shared memory per thread
  - $48K/32 = 1536$  threads
  - `occupancy=1`
- Example 2, 16K shared memory
  - Kernel uses 32 bytes of shared memory per thread
  - $16K/32 = 512$  threads
  - `occupancy=.3333`
- Don't use too much shared memory
- Choose L1/Shared config appropriately.



# Occupancy Limiter: Block Size

- Each SM can have up to 8 active blocks
- A small block size will limit the total number of threads

Block Size	Active Threads	Occupancy
32	256	.1666
64	512	.3333
128	1024	.6666
192	1536	1
256	2048 (1536)	1

- Avoid small block sizes, generally 128-256 threads is sufficient

# What Occupancy Do I Need?

- Depends on your problem...
  - Many find 66% is enough to saturate the bandwidth
- **Look at increasing occupancy only if the following are true!**
  - The kernel is bandwidth bound
  - The achieved bandwidth is significantly less than peak
- Instruction Level Parallelism (ILP) can have a greater effect than increasing occupancy
  - Vasily Volkov's GTC2010 talk "Better Performance at Lower Occupancy"
  - <http://nvidia.fullviewmedia.com/gtc2010/0922-a5-2238.html>



# Cuda Occupancy Calculator

- A tool to help you investigate occupancy
- [http://developer.download.nvidia.com/compute/cuda/4\\_0/sdk/docs/CUDA\\_Occupancy\\_Calculator.xls](http://developer.download.nvidia.com/compute/cuda/4_0/sdk/docs/CUDA_Occupancy_Calculator.xls)
- Demo: [CUDA Occupancy calculator.xls](#)



# Compute Profiler 4.0

- A useful profiling tool which can help you investigate occupancy, throughput, and bandwidth.
- Measures actual occupancy and thus may detect problems that shouldn't appear in theory
- Demo: [computeprof](#)

# Summary



- In order to achieve peak global memory bandwidth we need to have enough transactions in flight to hide latency
- We can increase the number of transactions by
  - Increasing occupancy
  - Increasing instruction level parallelism
- Occupancy can be limited by
  - Register usage
  - Shared memory usage
  - Block size
- Use the cuda occupancy calculator and the visual profiler to investigate memory bandwidth/occupancy



**Questions?**

